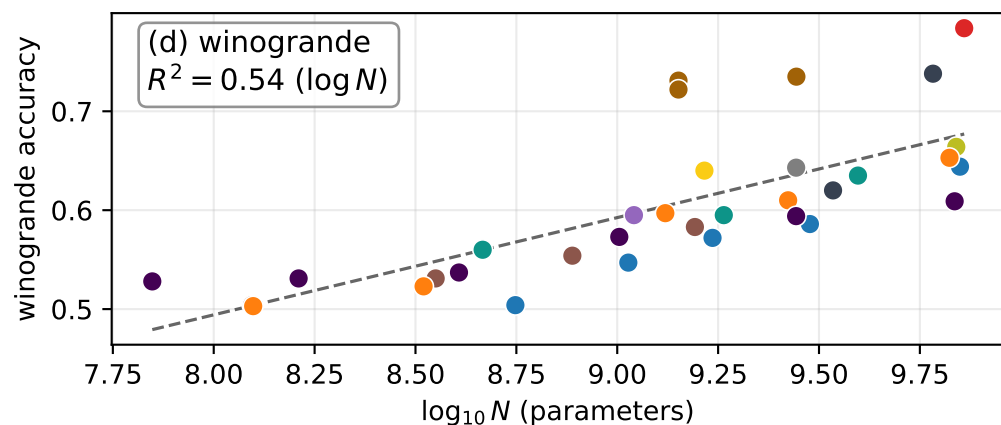
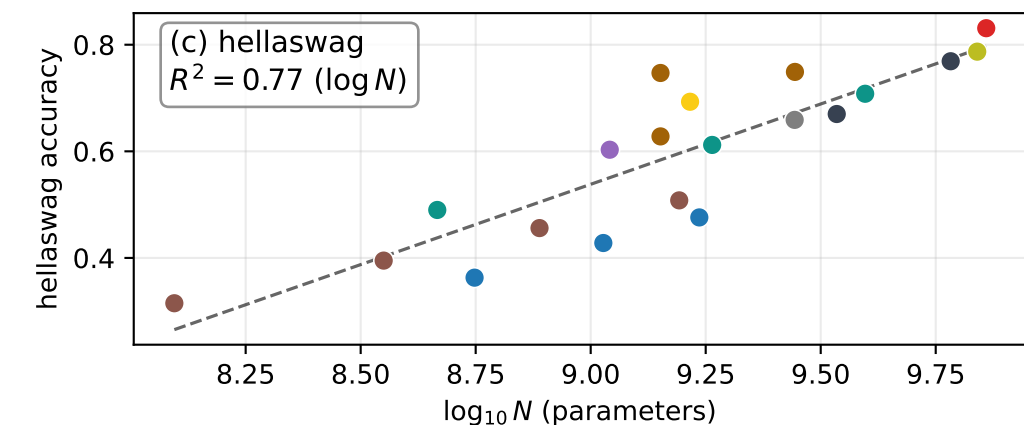
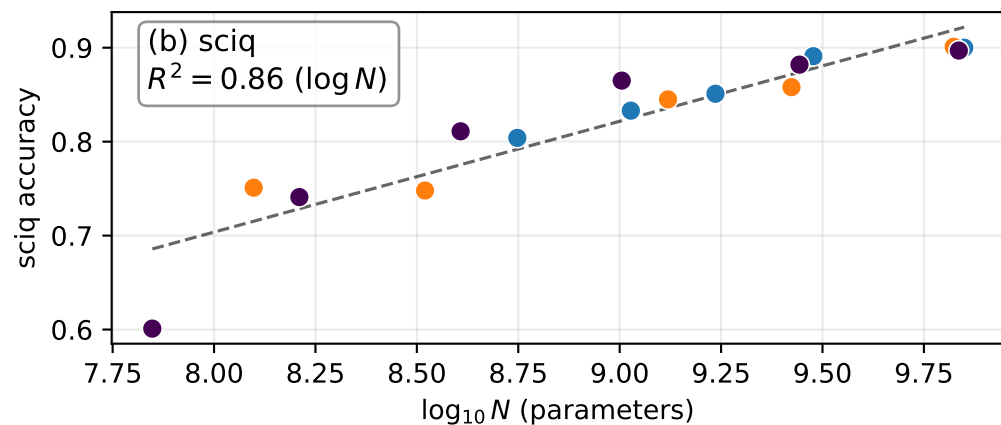
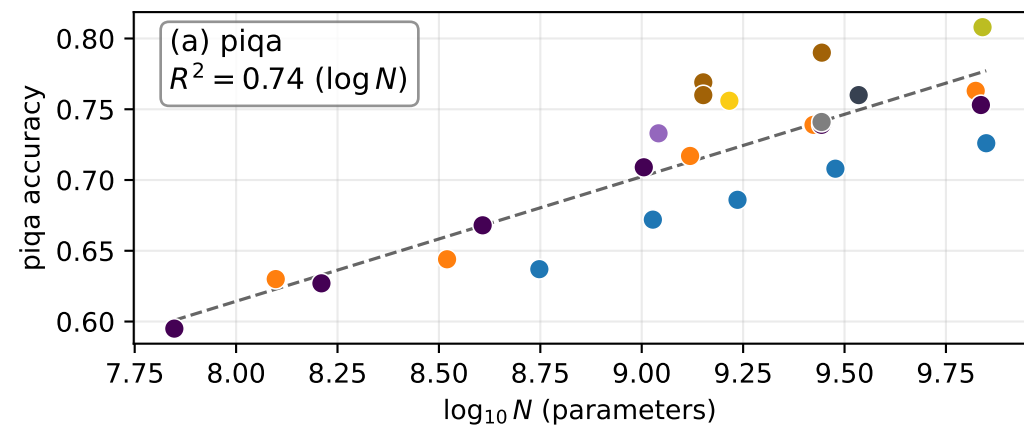


Figure 4. Spectral predictors of in-context-learning capability across 41 publicly released LMs

pythia opt falcon redpajama qwen mistral phi stablelm other
bloom gpt-2 tinyllama



(e) Spectral covariates raise R^2 over $\log N$ baseline (41 models, 8 benchmarks)

